

CSCI 6515 - Machine Learning for Big Data (Fall 2023)

Final Project

Group_ID: 7

Group Members:

1. Xiaoting_B00685239
2. Xuelian_B00977221
3. Siqi_B00976996

1. Dataset Information

Dataset Name: Credit_approval

Link to the Dataset:

<https://archive.ics.uci.edu/dataset/27/credit+approval>

Dataset Description:

This file concerns credit card applications. All attribute names and values have been changed to meaningless symbols to protect confidentiality of the data. This dataset is interesting because there is a good mix of attributes -- continuous, nominal with small numbers of values, and nominal with larger numbers of values.

2. Task Information

Task Goal: Predict the final decision of credit application.

Task Description:

The task is designed to predict the final decision on a credit application based on the applicant's personal and financial situation (e.g. income, marriage, debt, etc.). Four models were used for comparison of prediction performance, namely logistic regression, random forest, KNN and Bayesian model.

3. Task Implementation: Coding

3.1 Preprocessing

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: import ssl
import warnings
warnings.filterwarnings('ignore')
print(ssl.get_default_verify_paths())
```

DefaultVerifyPaths(cafile=None, capath=None, openssl_cafile_env='SSL_CERT_FILE', openssl_cafile='C:\\Program Files\\Common Files\\SSL\\cert.pem', openssl_capath_env='SSL_CERT_DIR', openssl_capath='C:\\Program Files\\Common Files\\SSL\\certs')

```
In [ ]: from ucimlrepo import fetch_ucirepo

# fetch dataset
credit_approval = fetch_ucirepo(id=27)

# data (as pandas dataframes)
X = credit_approval.data.features
y = credit_approval.data.targets
```

```
In [ ]: X = X.iloc[:, [14,13,12,11,10,9,8,7,6,5,4,3,2,1,0]]
X
```

```
Out[ ]:
```

		A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11	A12	A13	A14
0	b	30.83	0.000	u	g	w	v	1.25	t	t	1	f	g	202.0	
1	a	58.67	4.460	u	g	q	h	3.04	t	t	6	f	g	43.0	
2	a	24.50	0.500	u	g	q	h	1.50	t	f	0	f	g	280.0	
3	b	27.83	1.540	u	g	w	v	3.75	t	t	5	t	g	100.0	
4	b	20.17	5.625	u	g	w	v	1.71	t	f	0	f	s	120.0	
...	
685	b	21.08	10.085	y	p	e	h	1.25	f	f	0	f	g	260.0	
686	a	22.67	0.750	u	g	c	v	2.00	f	t	2	t	g	200.0	
687	a	25.25	13.500	y	p	ff	ff	2.00	f	t	1	t	g	200.0	
688	b	17.92	0.205	u	g	aa	v	0.04	f	f	0	f	g	280.0	
689	b	35.00	3.375	u	g	c	h	8.29	f	f	0	t	g	0.0	

690 rows x 15 columns

```
In [ ]: y
```

```
Out[ ]:
```

	A16
0	+
1	+
2	+
3	+
4	+
...	...
685	-
686	-
687	-
688	-
689	-

690 rows × 1 columns

```
In [ ]: df = X.merge(y,how='outer', left_index=True, right_index=True)
df
```

```
Out[ ]:
```

	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11	A12	A13	A14
0	b	30.83	0.000	u	g	w	v	1.25	t	t	1	f	g	202.0
1	a	58.67	4.460	u	g	q	h	3.04	t	t	6	f	g	43.0
2	a	24.50	0.500	u	g	q	h	1.50	t	f	0	f	g	280.0
3	b	27.83	1.540	u	g	w	v	3.75	t	t	5	t	g	100.0
4	b	20.17	5.625	u	g	w	v	1.71	t	f	0	f	s	120.0
...
685	b	21.08	10.085	y	p	e	h	1.25	f	f	0	f	g	260.0
686	a	22.67	0.750	u	g	c	v	2.00	f	t	2	t	g	200.0
687	a	25.25	13.500	y	p	ff	ff	2.00	f	t	1	t	g	200.0
688	b	17.92	0.205	u	g	aa	v	0.04	f	f	0	f	g	280.0
689	b	35.00	3.375	u	g	c	h	8.29	f	f	0	t	g	0.0

690 rows × 16 columns

```
In [ ]: df = df.rename(columns={"A1":"Gender","A2":"Age","A3":"Debt","A4":"Married",
                                "A7":"Ethnicity","A8":"YearsEmployed","A9":"PriorCriminalRecord",
                                "A12":"DriverLicense","A13":"Citizen","A14":"ZipCode"})
df
```

```
Out [ ]:
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	Y
0	b	30.83	0.000	u	g	w	v	
1	a	58.67	4.460	u	g	q	h	
2	a	24.50	0.500	u	g	q	h	
3	b	27.83	1.540	u	g	w	v	
4	b	20.17	5.625	u	g	w	v	
...
685	b	21.08	10.085	y	p	e	h	
686	a	22.67	0.750	u	g	c	v	
687	a	25.25	13.500	y	p	ff	ff	
688	b	17.92	0.205	u	g	aa	v	
689	b	35.00	3.375	u	g	c	h	

690 rows x 16 columns

```
In [ ]: df.isnull().sum()
```

```
Out [ ]: Gender          12
Age              12
Debt              0
Married          6
BankCustomer     6
EducationLevel   9
Ethnicity        9
YearsEmployed    0
PriorDefault     0
Employed         0
CreditScore     0
DriverLicense    0
Citizen          0
ZipCode         13
Income           0
ApprovalStatus   0
dtype: int64
```

```
In [ ]: df = df.dropna(axis=0)
df = df.drop_duplicates()
df
```

```
Out [ ]:
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	Y
0	b	30.83	0.000	u	g	w	v	
1	a	58.67	4.460	u	g	q	h	
2	a	24.50	0.500	u	g	q	h	
3	b	27.83	1.540	u	g	w	v	
4	b	20.17	5.625	u	g	w	v	
...
685	b	21.08	10.085	y	p	e	h	
686	a	22.67	0.750	u	g	c	v	
687	a	25.25	13.500	y	p	ff	ff	
688	b	17.92	0.205	u	g	aa	v	
689	b	35.00	3.375	u	g	c	h	

653 rows x 16 columns

```
In [ ]: df.isnull().sum()
```

```
Out [ ]: Gender          0
Age          0
Debt         0
Married      0
BankCustomer 0
EducationLevel 0
Ethnicity    0
YearsEmployed 0
PriorDefault 0
Employed     0
CreditScore  0
DriverLicense 0
Citizen      0
ZipCode      0
Income       0
ApprovalStatus 0
dtype: int64
```

Encode and Normalization

```
In [ ]: df['ApprovalStatus'] = df['ApprovalStatus'].map({'+' : 1, '-' : 0})
df
```

```
Out[ ]:
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	Y
0	b	30.83	0.000	u	g	w	v	
1	a	58.67	4.460	u	g	q	h	
2	a	24.50	0.500	u	g	q	h	
3	b	27.83	1.540	u	g	w	v	
4	b	20.17	5.625	u	g	w	v	
...	
685	b	21.08	10.085	y	p	e	h	
686	a	22.67	0.750	u	g	c	v	
687	a	25.25	13.500	y	p	ff	ff	
688	b	17.92	0.205	u	g	aa	v	
689	b	35.00	3.375	u	g	c	h	

653 rows x 16 columns

```
In [ ]: from sklearn.preprocessing import LabelEncoder
#Apply label encoding on all columns with type of object
df[list(df.columns[df.dtypes == object])] = df[list(df.columns[df.dtypes == object])]
df
```

```
Out[ ]:
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity	Y
0	1	30.83	0.000	1	0	12	7	
1	0	58.67	4.460	1	0	10	3	
2	0	24.50	0.500	1	0	10	3	
3	1	27.83	1.540	1	0	12	7	
4	1	20.17	5.625	1	0	12	7	
...	
685	1	21.08	10.085	2	2	4	3	
686	0	22.67	0.750	1	0	1	7	
687	0	25.25	13.500	2	2	5	2	
688	1	17.92	0.205	1	0	0	7	
689	1	35.00	3.375	1	0	1	3	

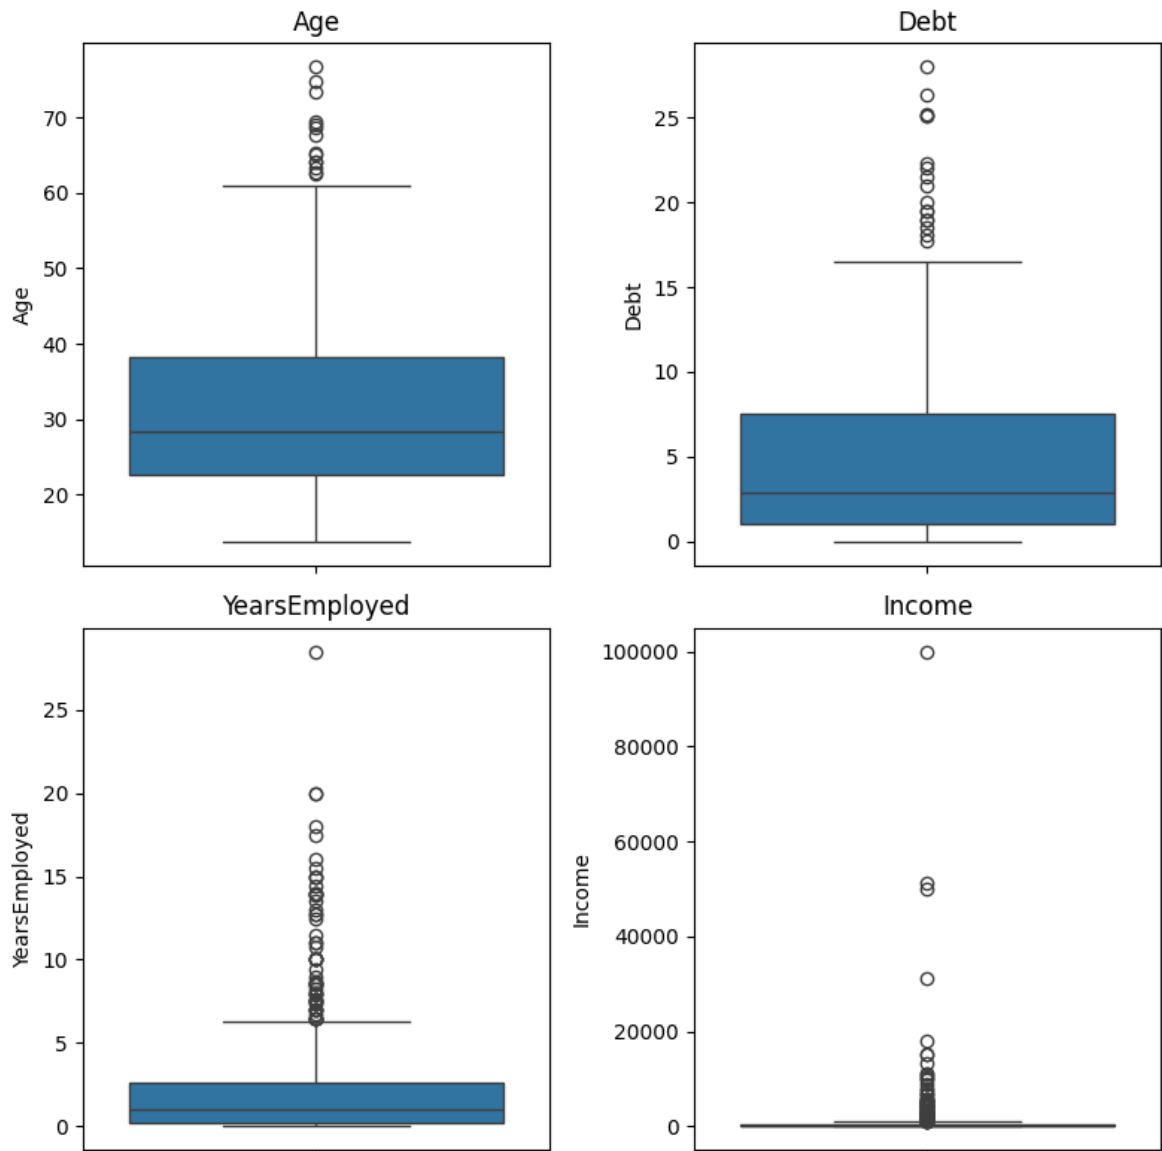
653 rows x 16 columns

```
In [ ]: # Plotting boxplots for continuous columns to identify outliers
plt.figure(figsize=(8, 8))
columns = ['Age', 'Debt', 'YearsEmployed', 'Income']

for i, col in enumerate(columns, 1):
    plt.subplot(2, 2, i)
    #sns.boxplot(y=train_data[col])
    #sns.boxplot(data = df , x = 'ApprovalStatus' , y = col)
```

```
sns.boxplot(data = df , y = col)
plt.title(col)

plt.tight_layout()
plt.show()
```



```
In [ ]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[columns] = scaler.fit_transform(df[columns])
df.head()
```


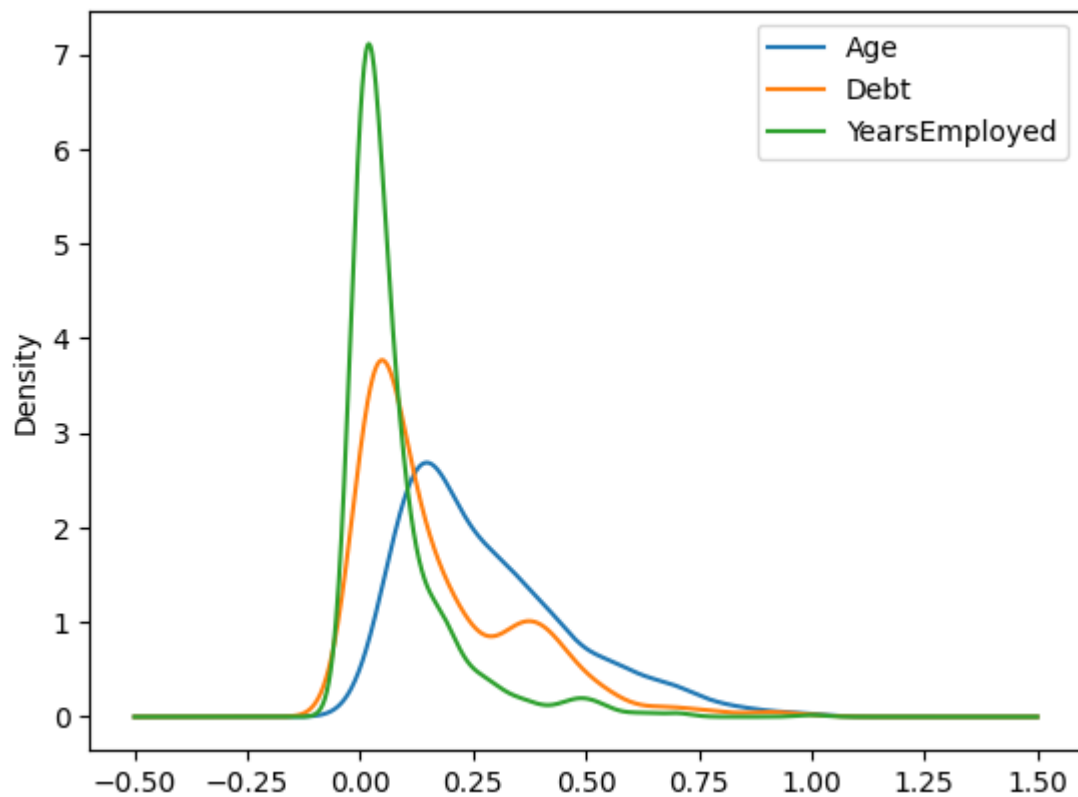
```
Out[ ]:
```

	Gender	Age	Debt	Married	BankCustomer	EducationLevel	Ethnicity
0	1	0.271111	0.000000	1	0	12	7
1	0	0.713016	0.159286	1	0	10	3
2	0	0.170635	0.017857	1	0	10	3
3	1	0.223492	0.055000	1	0	12	7
4	1	0.101905	0.200893	1	0	12	7


Check the distribution of variables

```
In [ ]: df.iloc[:, [1,2,7]].plot.kde()  
df.iloc[:, [14]].plot.kde()
```

Out[]: <Axes: ylabel='Density'>



```
In [ ]: sns.heatmap(df.corr(), annot=True, cmap='coolwarm')  
plt.gcf().set_size_inches(15, 15)  
plt.show()
```

3.2 Model development and evaluation

3.2.1 Logistic Regression

```
In [ ]: from sklearn.model_selection import train_test_split, cross_val_score
        from sklearn.metrics import classification_report, accuracy_score

        X = df.iloc[:, 0:15]
        y = df.iloc[:, -1:]

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

3.2.1.1 Hyperparameter tune

```
In [ ]: from sklearn.model_selection import GridSearchCV
        from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import classification_report, accuracy_score
        # define param grid
        param_grid = {
```

```

'C': [0.001, 0.01, 0.1, 1, 10],
'penalty': ['l1', 'l2', 'elasticnet'],
'tol': [1e-4, 1e-3, 1e-2],
'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
}

# build model
logistic = LogisticRegression(max_iter=1000)

# GridSearchCV
grid_search = GridSearchCV(logistic, param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

# get best params and score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

# use best params
best_logistic = LogisticRegression(**best_params, max_iter=1000)
best_logistic.fit(X_train, y_train)
y_pred_best = best_logistic.predict(X_test)
accuracy_best = accuracy_score(y_test, y_pred_best)
class_report = classification_report(y_test, y_pred_best)

print("Best Parameters:", best_params)
print("Best Cross-Validation Score:", best_score)
print("Test Set Accuracy:", accuracy_best)
print("Classification Report:\n", class_report)

```

Best Parameters: {'C': 1, 'penalty': 'l1', 'solver': 'liblinear', 'tol': 0.001}

Best Cross-Validation Score: 0.8642618251313904

Test Set Accuracy: 0.8877551020408163

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.88	0.90	112
1	0.84	0.90	0.87	84
accuracy			0.89	196
macro avg	0.88	0.89	0.89	196
weighted avg	0.89	0.89	0.89	196

3.2.1.2 Model evaluation

```

In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# get result
results = pd.DataFrame(grid_search.cv_results_)

# choose parameters
param_1 = 'param_C'
param_2 = 'param_solver'

# plot


```

```

pivot_table = results.pivot_table(values='mean_test_score',
                                   index=[param_1],
                                   columns=[param_2])

plt.figure(figsize=(10, 6))
sns.heatmap(pivot_table, annot=True, cmap='YlGnBu')
plt.title('Grid Search Scores')
plt.xlabel('Solver')
plt.ylabel('C (Regularization Strength)')
plt.show()

```



3.2.2 KNN Classifier

```

In [ ]: from sklearn.model_selection import train_test_split, cross_val_score
        from sklearn.metrics import classification_report, accuracy_score
        from sklearn.neighbors import KNeighborsClassifier

X = df.iloc[:, 0:15]
y = df.iloc[:, -1:]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

KNN = KNeighborsClassifier(n_neighbors=10)
KNN.fit(X_train.to_numpy(), y_train.to_numpy())
y_pred = KNN.predict(X_test.to_numpy())

cm = classification_report(y_test.to_numpy(), y_pred)
print(cm)

```

	precision	recall	f1-score	support
0	0.64	0.86	0.73	109
1	0.69	0.39	0.50	87
accuracy			0.65	196
macro avg	0.67	0.63	0.62	196
weighted avg	0.66	0.65	0.63	196

3.2.2.1 Hyperparameter tune


```
In [ ]: k_range = range(1,50)
cv_range = [2,3,4,5,6,7,8,9,10]
k_scores = []
cv_scores = []

#for i in cv_range:
for k in k_range:
    KNN = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(KNN, X_train.to_numpy(), y_train.to_numpy(),
                             k_scores.append(scores.mean()))
    #cv_scores[i,:] = k_scores
print(k_scores)
print(k_scores.index(max(k_scores)))
```

```
[0.6761108456760629, 0.6432632584806497, 0.7090301003344482, 0.68927376970
85524, 0.715551839464883, 0.6827281414237938, 0.6958910654562829, 0.678332
5370281892, 0.6870520783564262, 0.6914715719063544, 0.6740324892498805, 0.
6740324892498805, 0.6696368848542761, 0.6586478738652651, 0.65210224558050
64, 0.6543000477783086, 0.678356426182513, 0.6717391304347826, 0.665193502
1500238, 0.6717869087434305, 0.6564739608217869, 0.6498805542283803, 0.641
0654562828475, 0.6498805542283803, 0.6345676063067367, 0.6433349259436215,
0.6455566172957478, 0.6477305303392259, 0.6499283325370282, 0.652078356426
1825, 0.6564739608217869, 0.6674390826564739, 0.6478260869565218, 0.65, 0.
6521739130434783, 0.6631151457238414, 0.6455566172957478, 0.64335881509794
56, 0.6368131868131869, 0.6412087912087913, 0.6498805542283803, 0.65434782
60869566, 0.6629956999522217, 0.6695652173913043, 0.6673674151935021, 0.66
08695652173913, 0.6476588628762541, 0.6564261825131391, 0.641137123745819
3]
4
```

```
In [ ]: plt.plot(k_range, k_scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')
```

```
Out[ ]: Text(0, 0.5, 'Cross-Validated Accuracy')
```



3.2.2.2 Model evaluation

```
In [ ]: KNN = KNeighborsClassifier(n_neighbors=k_range[k_scores.index(max(k_score
KNN.fit(X_train.to_numpy(), y_train.to_numpy())
y_pred = KNN.predict(X_test.to_numpy())

cm = classification_report(y_test.to_numpy(),y_pred)
print(cm)
```

	precision	recall	f1-score	support
0	0.66	0.78	0.72	109
1	0.65	0.51	0.57	87
accuracy			0.66	196
macro avg	0.66	0.64	0.64	196
weighted avg	0.66	0.66	0.65	196

3.2.3 Naive Bayes Classifier

```
In [ ]: from sklearn.naive_bayes import GaussianNB
GaussNB = GaussianNB()
GaussNB.fit(X_train, y_train)
y_pred = GaussNB.predict(X_test)
cm = classification_report(y_test,y_pred)
print(cm)
```

	precision	recall	f1-score	support
0	0.77	0.97	0.86	109
1	0.95	0.63	0.76	87
accuracy			0.82	196
macro avg	0.86	0.80	0.81	196
weighted avg	0.85	0.82	0.81	196

```
In [ ]: from sklearn.naive_bayes import BernoulliNB
BernouNB = BernoulliNB(force_alpha=True)
BernouNB.fit(X_train, y_train)
y_pred = BernouNB.predict(X_test)
cm = classification_report(y_test, y_pred)
print(cm)
```

	precision	recall	f1-score	support
0	0.80	0.90	0.84	109
1	0.85	0.71	0.77	87
accuracy			0.82	196
macro avg	0.82	0.81	0.81	196
weighted avg	0.82	0.82	0.81	196

```
In [ ]: from sklearn.naive_bayes import MultinomialNB
MultiNomNB = MultinomialNB()
MultiNomNB.fit(X_train, y_train)
y_pred = MultiNomNB.predict(X_test)
cm = classification_report(y_test, y_pred)
print(cm)
```

	precision	recall	f1-score	support
0	0.73	0.76	0.74	109
1	0.68	0.64	0.66	87
accuracy			0.71	196
macro avg	0.71	0.70	0.70	196
weighted avg	0.71	0.71	0.71	196

3.2.4 Random Forest Classifier

```
In [ ]: from sklearn.model_selection import train_test_split, RandomizedSearchCV,
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [ ]: X = df.iloc[:,0:15]
y = df.iloc[:,15:]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.3)
```

3.2.4.3 Hyperparameter tune

```
In [ ]: rf_classifier = RandomForestClassifier(random_state=42)
```

```
In [ ]: param_dist = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    # 'max_features': ['auto', 'sqrt', 'log2']
}
```

```
In [ ]: random_search = RandomizedSearchCV(rf_classifier, param_distributions=param_dist,
    random_search.fit(X_train, y_train)
```

```
Out[ ]: RandomizedSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier
```

```
In [ ]: print("Best Hyperparameters_RandomSearch:", random_search.best_params_)
```

Best Hyperparameters_RandomSearch: {'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 10}

```
In [ ]: best_rf_model_Random = random_search.best_estimator_
    y_pred = best_rf_model_Random.predict(X_val)
```

```
In [ ]: accuracy = accuracy_score(y_val, y_pred)
    print("Random Search Validation Accuracy:", accuracy)
```

Random Search Validation Accuracy: 0.8913043478260869

```
In [ ]: confusion_mat = confusion_matrix(y_val, y_pred)
    print("Random Search Confusion Matrix:\n", confusion_mat)
```

Random Search Confusion Matrix:
[[47 5]
[5 35]]

```
In [ ]: class_report = classification_report(y_val, y_pred)
    print("Random Search Classification Report:\n", class_report)
```

Random Search Classification Report:

	precision	recall	f1-score	support
0	0.90	0.90	0.90	52
1	0.88	0.88	0.88	40
accuracy			0.89	92
macro avg	0.89	0.89	0.89	92
weighted avg	0.89	0.89	0.89	92

```
In [ ]: param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    # 'max_features': ['auto', 'sqrt', 'log2']
}
```

```
In [ ]: grid_search = GridSearchCV(rf_classifier, param_grid, cv=5, scoring='accuracy')
        grid_search.fit(X_train, y_train)
```

```
Out [ ]:
└─ GridSearchCV
  └─ estimator: RandomForestClassifier
    └─ RandomForestClassifier
```

```
In [ ]: print("Best Hyperparameters_GridSearch:", grid_search.best_params_)
```

```
Best Hyperparameters_GridSearch: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 100}
```

****The processing time of Grid Search is much slower than Random Search. The different results of these two kinds of hyperparameter tuning methods are:**

Random Search: "min_samples_split": 2, "n_estimators": 50,

Grid Search: "min_samples_split": 10, "n_estimators": 150

Compare the performance for the validation set: (Random Search already done) **

```
In [ ]: best_rf_model_Grid = grid_search.best_estimator_
        y_pred = best_rf_model_Grid.predict(X_val)
```

```
In [ ]: accuracy = accuracy_score(y_val, y_pred)
        print("Grid Search Validation Accuracy:", accuracy)
```

```
Grid Search Validation Accuracy: 0.8695652173913043
```

```
In [ ]: confusion_mat = confusion_matrix(y_val, y_pred)
        print("Grid Search Confusion Matrix:\n", confusion_mat)
```

```
Grid Search Confusion Matrix:
[[47  5]
 [ 7 33]]
```

```
In [ ]: class_report = classification_report(y_val, y_pred)
        print("Grid Search Classification Report:\n", class_report)
```

```
Grid Search Classification Report:
              precision    recall  f1-score   support

     0           0.87         0.90         0.89         52
     1           0.87         0.82         0.85         40

   accuracy              0.87              92
  macro avg           0.87         0.86         0.87              92
 weighted avg           0.87         0.87         0.87              92
```

3.2.4.3 Model evaluation

The performances from the random grid is better. Just use the hyperparameter from random grid to train the test dataset

```
In [ ]: best_rf_model_Random = random_search.best_estimator_
        y_test_pred = best_rf_model_Random.predict(X_test)
```



```
In [ ]: accuracy = accuracy_score(y_test, y_test_pred)
        print("Test Accuracy:", accuracy)
```

Test Accuracy: 0.8775510204081632

```
In [ ]: confusion_mat = confusion_matrix(y_test, y_test_pred)
        print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[105  9]
 [ 15 67]]
```

```
In [ ]: class_report = classification_report(y_test, y_test_pred)
        print("Classification Report:\n", class_report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.92	0.90	114
1	0.88	0.82	0.85	82
accuracy			0.88	196
macro avg	0.88	0.87	0.87	196
weighted avg	0.88	0.88	0.88	196